

In [31]:

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
```

Type *Markdown* and LaTeX:  $\alpha^2$

In [32]:

```
data1=pd.read_csv("train.csv")
data2=pd.read_csv("test.csv")
data = pd.concat([data1,data2],axis=0)
```


In [33]:

```
data.head(5)
```

Out[33]:

	Unnamed: 0	id	Gender	Customer Type	Age	Type of Travel	Class	Flight Distance	Inflight wifi service	Departure/ time con
0	0	70172	Male	Loyal Customer	13	Personal Travel	Eco Plus	460	3	
1	1	5047	Male	disloyal Customer	25	Business travel	Business	235	3	
2	2	110028	Female	Loyal Customer	26	Business travel	Business	1142	2	
3	3	24026	Female	Loyal Customer	25	Business travel	Business	562	2	
4	4	119299	Male	Loyal Customer	61	Business travel	Business	214	3	

5 rows × 25 columns



```
In [34]: print(data.info())
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 129880 entries, 0 to 25975
Data columns (total 25 columns):
#   Column                                          Non-Null Count  Dtype
---  -
0   Unnamed: 0                                     129880 non-null  int64
1   id                                              129880 non-null  int64
2   Gender                                         129880 non-null  object
3   Customer Type                                129880 non-null  object
4   Age                                            129880 non-null  int64
5   Type of Travel                               129880 non-null  object
6   Class                                          129880 non-null  object
7   Flight Distance                              129880 non-null  int64
8   Inflight wifi service                        129880 non-null  int64
9   Departure/Arrival time convenient            129880 non-null  int64
10  Ease of Online booking                       129880 non-null  int64
11  Gate location                                129880 non-null  int64
12  Food and drink                               129880 non-null  int64
13  Online boarding                              129880 non-null  int64
14  Seat comfort                                 129880 non-null  int64
15  Inflight entertainment                      129880 non-null  int64
16  On-board service                             129880 non-null  int64
17  Leg room service                             129880 non-null  int64
18  Baggage handling                             129880 non-null  int64
19  Checkin service                             129880 non-null  int64
20  Inflight service                             129880 non-null  int64
21  Cleanliness                                  129880 non-null  int64
22  Departure Delay in Minutes                   129880 non-null  int64
23  Arrival Delay in Minutes                     129487 non-null  float64
24  satisfaction                                  129880 non-null  object
dtypes: float64(1), int64(19), object(5)
memory usage: 25.8+ MB
None
```

```
In [35]: del data["Unnamed: 0"]
del data["id"]
```

```
In [36]: print(data.isnull().sum())
```

Gender	0
Customer Type	0
Age	0
Type of Travel	0
Class	0
Flight Distance	0
Inflight wifi service	0
Departure/Arrival time convenient	0
Ease of Online booking	0
Gate location	0
Food and drink	0
Online boarding	0
Seat comfort	0
Inflight entertainment	0
On-board service	0
Leg room service	0
Baggage handling	0
Checkin service	0
Inflight service	0
Cleanliness	0
Departure Delay in Minutes	0
Arrival Delay in Minutes	393
satisfaction	0
dtype:	int64

```
In [37]: from sklearn.impute import SimpleImputer
imputer_int=SimpleImputer(missing_values=np.nan)
data['Arrival Delay in Minutes'] = imputer_int.fit_transform(data[['Arrival De
```

```
In [38]: data.isnull().sum()
```

```
Out[38]: Gender                                0
Customer Type                                0
Age                                            0
Type of Travel                               0
Class                                          0
Flight Distance                              0
Inflight wifi service                        0
Departure/Arrival time convenient            0
Ease of Online booking                       0
Gate location                                0
Food and drink                              0
Online boarding                             0
Seat comfort                                0
Inflight entertainment                       0
On-board service                            0
Leg room service                            0
Baggage handling                            0
Checkin service                             0
Inflight service                            0
Cleanliness                                 0
Departure Delay in Minutes                   0
Arrival Delay in Minutes                     0
satisfaction                                0
dtype: int64
```

```
In [39]: from sklearn.preprocessing import LabelEncoder
LE=LabelEncoder()
data["Gender"]=LE.fit_transform(data["Gender"])
data["Customer Type"]=LE.fit_transform(data["Customer Type"])
data["Type of Travel"]=LE.fit_transform(data["Type of Travel"])
data["Class"]=LE.fit_transform(data["Class"])
data["satisfaction"]=LE.fit_transform(data["satisfaction"])
```

In [40]: data.info()

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 129880 entries, 0 to 25975
Data columns (total 23 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Gender                                129880 non-null  int32
1   Customer Type                        129880 non-null  int32
2   Age                                  129880 non-null  int64
3   Type of Travel                       129880 non-null  int32
4   Class                                129880 non-null  int32
5   Flight Distance                      129880 non-null  int64
6   Inflight wifi service                129880 non-null  int64
7   Departure/Arrival time convenient    129880 non-null  int64
8   Ease of Online booking                129880 non-null  int64
9   Gate location                        129880 non-null  int64
10  Food and drink                       129880 non-null  int64
11  Online boarding                      129880 non-null  int64
12  Seat comfort                         129880 non-null  int64
13  Inflight entertainment               129880 non-null  int64
14  On-board service                    129880 non-null  int64
15  Leg room service                    129880 non-null  int64
16  Baggage handling                    129880 non-null  int64
17  Checkin service                     129880 non-null  int64
18  Inflight service                    129880 non-null  int64
19  Cleanliness                         129880 non-null  int64
20  Departure Delay in Minutes           129880 non-null  int64
21  Arrival Delay in Minutes             129880 non-null  float64
22  satisfaction                         129880 non-null  int32
dtypes: float64(1), int32(5), int64(17)
memory usage: 21.3 MB
```

```
In [41]: IndepVar = []
for col in data.columns:
    if col != 'satisfaction':
        IndepVar.append(col)

TargetVar = 'satisfaction'

x = data[IndepVar]
y = data[TargetVar]
```


In [42]: x.columns

```
Out[42]: Index(['Gender', 'Customer Type', 'Age', 'Type of Travel', 'Class',
               'Flight Distance', 'Inflight wifi service',
               'Departure/Arrival time convenient', 'Ease of Online booking',
               'Gate location', 'Food and drink', 'Online boarding', 'Seat comfort',
               'Inflight entertainment', 'On-board service', 'Leg room service',
               'Baggage handling', 'Checkin service', 'Inflight service',
               'Cleanliness', 'Departure Delay in Minutes',
               'Arrival Delay in Minutes'],
              dtype='object')
```

```
In [44]: corr = data.corr()
```

```
In [45]: sns.heatmap(corr)
```

```
Out[45]: <AxesSubplot:>
```



```
In [46]: from sklearn.model_selection import train_test_split
x_train,x_test, y_train, y_test=train_test_split(x,y,test_size=0.33,random_sta
```

```
In [47]: from sklearn.preprocessing import StandardScaler
sc=StandardScaler()
x_train=sc.fit_transform(x_train)
x_test=sc.transform (x_test)
```

```
In [48]: print(x_train.shape,x_test.shape)
```

```
(87019, 22) (42861, 22)
```

## Selecting Meaningful Features

```
In [19]: from sklearn.feature_selection import SelectKBest,chi2
best = SelectKBest(score_func=chi2,k=10)
fit = best.fit(x,y)
scores = pd.DataFrame(fit.scores_)
colNames = pd.DataFrame(x.columns)
features = pd.concat([colNames,scores],axis=1)
features.columns = ['attribute','score']
imp_features=features.nlargest(10,'score')
print(imp_features)
```

	attribute	score
5	Flight Distance	9.645859e+06
21	Arrival Delay in Minutes	4.298388e+04
20	Departure Delay in Minutes	3.293872e+04
11	Online boarding	1.834034e+04
3	Type of Travel	1.815163e+04
4	Class	1.696271e+04
2	Age	1.353949e+04
13	Inflight entertainment	1.091611e+04
12	Seat comfort	7.993086e+03
6	Inflight wifi service	6.758345e+03

```
In [20]: c = ['Flight Distance','Arrival Delay in Minutes','Departure Delay in Minutes'
'Inflight entertainment','Seat comfort','Inflight wifi service']
x1=data.loc[:,c]
print(x1.head(1))
```

	Flight Distance	Arrival Delay in Minutes	Departure Delay in Minutes	\
0	460	18.0	25	

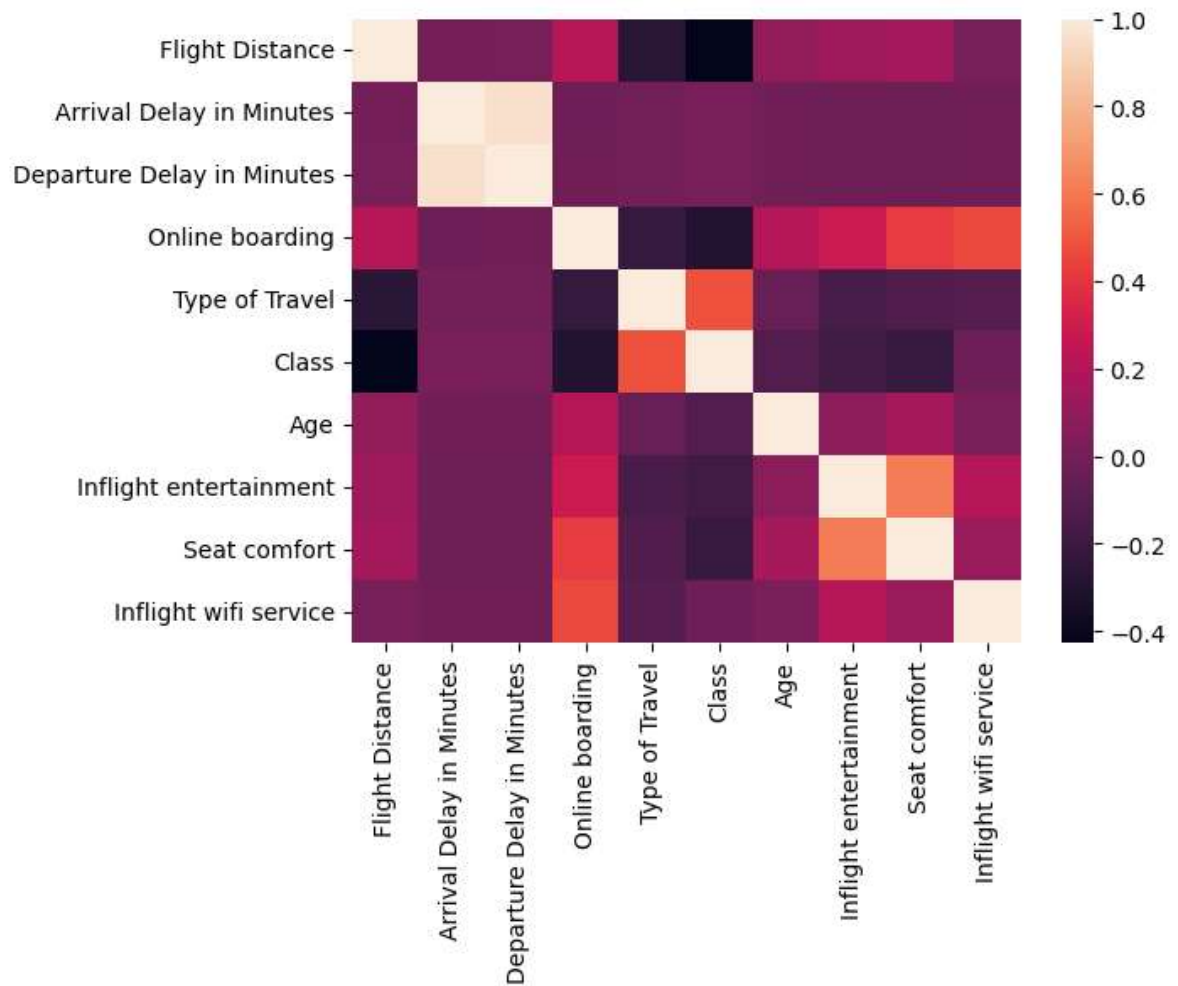
	Online boarding	Type of Travel	Class	Age	Inflight entertainment	\
0	3	1	2	13	5	

	Seat comfort	Inflight wifi service
0	5	3

```
In [21]: x1corr = x1.corr()
sns.heatmap(x1corr)
```

Out[21]: <AxesSubplot:>



```
In [22]: from sklearn.model_selection import train_test_split
x_train,x_test, y_train, y_test=train_test_split(x1,y,test_size=0.33,random_st
```

```
In [23]: from sklearn.preprocessing import StandardScaler
sc=StandardScaler()
x_train=sc.fit_transform(x_train)
x_test=sc.transform (x_test)
```

```
In [24]: print(x_train.shape,x_test.shape)
```

(87019, 10) (42861, 10)

In [ ]:

## PCA



```
In [25]: from sklearn.decomposition import PCA
pca=PCA(n_components = 2)
x_r=pca.fit(x).transform (x)
```

```
In [26]: print(x_r)

[[-730.35390612    9.38687236]
 [-955.33902476   -16.04965605]
 [-48.33502762   -21.0399969 ]
 ...
 [-362.3515567    -20.99952443]
 [-63.35407309   -21.00030257]
 [-926.31430049   -21.0677217 ]]
```

```
In [ ]:
```

## Common Techniques to identify overfitting and underfitting:

1.learning curves 2.cross-validation 3.Regularization Techniques

```
In [27]: #using Regularization Technique
from sklearn.linear_model import LogisticRegression
# Initialize logistic regression with L2 regularization
clf = LogisticRegression(penalty='l2', C=1.0)
# Train the model and evaluate
clf.fit(x_train, y_train)
train_accuracy = clf.score(x_train, y_train)
test_accuracy = clf.score(x_test, y_test)
print("Training accuracy:", train_accuracy)
print("Testing accuracy:", test_accuracy)
```

```
Training accuracy: 0.8411611257311622
Testing accuracy: 0.8430274608618558
```

```
In [28]: from sklearn.model_selection import cross_val_score
# Perform cross-validation
scores = cross_val_score(clf, x, y, cv=5) # clf is your trained model, X is y
print("Cross-validation scores:", scores)
print("Mean accuracy:", np.mean(scores))
```

C:\Users\srira\anaconda3\lib\site-packages\sklearn\linear\_model\\_logistic.py:  
814: ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:  
<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)  
Please also refer to the documentation for alternative solver options:  
[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression) ([https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression))  
n\_iter\_i = \_check\_optimize\_result(  
C:\Users\srira\anaconda3\lib\site-packages\sklearn\linear\_model\\_logistic.py:  
814: ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:  
<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)  
Please also refer to the documentation for alternative solver options:  
[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression) ([https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression))  
n\_iter\_i = \_check\_optimize\_result(  
C:\Users\srira\anaconda3\lib\site-packages\sklearn\linear\_model\\_logistic.py:  
814: ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:  
<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)  
Please also refer to the documentation for alternative solver options:  
[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression) ([https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression))  
n\_iter\_i = \_check\_optimize\_result(  
C:\Users\srira\anaconda3\lib\site-packages\sklearn\linear\_model\\_logistic.py:  
814: ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:  
<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)  
Please also refer to the documentation for alternative solver options:  
[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression) ([https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression))  
n\_iter\_i = \_check\_optimize\_result(  
C:\Users\srira\anaconda3\lib\site-packages\sklearn\linear\_model\\_logistic.py:  
814: ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Cross-validation scores: [0.82884201 0.81817832 0.82025716 0.81729289 0.80031568]

Mean accuracy: 0.8169772097320603

C:\Users\sraira\anaconda3\lib\site-packages\sklearn\linear\_model\\_logistic.py: 814: ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:  
<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression) ([https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression))

```
n_iter_i = _check_optimize_result(
```

## Applying different algorithms

### LogisticRegression

```
In [49]: from sklearn.linear_model import LogisticRegression
ModellR = LogisticRegression()

ModellR.fit(x_train, y_train)

y_pred = ModellR.predict(x_test)

y_pred_prob = ModellR.predict_proba(x_test)

print('Model Name: ',ModellR )

from sklearn.metrics import confusion_matrix
actual = y_test
predicted = y_pred
matrix = confusion_matrix(actual,predicted, labels=[1,0],sample_weight=None, r
print('Confusion matrix : \n', matrix)
```

Model Name: LogisticRegression()

Confusion matrix :

```
[[15534  3094]
 [ 2296 21937]]
```

```
In [52]: tp, fn, fp, tn = confusion_matrix(actual,predicted,labels=[1,0]).reshape(-1)
print('Outcome values : tp:{} fn:{} fp:{} tn:{}\n'.format(tp, fn, fp, tn) )

sensitivity = round(tp/(tp+fn), 3)
specificity = round(tn/(tn+fp), 3)
accuracy = round((tp+tn)/(tp+fp+tn+fn), 3)
precision = round(tp/(tp+fp), 3)
f1Score = round((2*tp/(2*tp + fp + fn)), 3)

from math import sqrt

print('Accuracy :', round(accuracy*100, 2),'%')
print('Precision :', round(precision*100, 2),'%')
print('Recall :', round(sensitivity*100,2), '%')
print('F1 Score :', f1Score)
print('Specificity or True Negative Rate :', round(specificity*100,2), '%' )
```

Outcome values : tp:15534 fn:3094 fp:2296 tn:21937

Accuracy : 87.4 %  
Precision : 87.1 %  
Recall : 83.4 %  
F1 Score : 0.852  
Specificity or True Negative Rate : 90.5 %

## Random Forest

```
In [53]: from sklearn.ensemble import RandomForestClassifier
ModelRF = RandomForestClassifier()
ModelLR.fit(x_train, y_train)

y_pred = ModelLR.predict(x_test)

y_pred_prob = ModelLR.predict_proba(x_test)

print('Model Name: ',ModelLR )

from sklearn.metrics import confusion_matrix
actual = y_test
predicted = y_pred
matrix = confusion_matrix(actual,predicted, labels=[1,0],sample_weight=None, r
print('Confusion matrix : \n', matrix)
```

Model Name: LogisticRegression()  
Confusion matrix :  
[[15534 3094]  
[ 2296 21937]]

```
In [54]: tp, fn, fp, tn = confusion_matrix(actual,predicted,labels=[1,0]).reshape(-1)
print('Outcome values : tp:{} fn:{} fp:{} tn:{}\n'.format(tp, fn, fp, tn) )

sensitivity = round(tp/(tp+fn), 3)
specificity = round(tn/(tn+fp), 3)
accuracy = round((tp+tn)/(tp+fp+tn+fn), 3)
precision = round(tp/(tp+fp), 3)
f1Score = round((2*tp/(2*tp + fp + fn)), 3)

from math import sqrt

print('Accuracy :', round(accuracy*100, 2),'%')
print('Precision :', round(precision*100, 2),'%')
print('Recall :', round(sensitivity*100,2), '%')
print('F1 Score :', f1Score)
print('Specificity or True Negative Rate :', round(specificity*100,2), '%' )
```

Outcome values : tp:15534 fn:3094 fp:2296 tn:21937

Accuracy : 87.4 %  
Precision : 87.1 %  
Recall : 83.4 %  
F1 Score : 0.852  
Specificity or True Negative Rate : 90.5 %

## KNeighbourClassifier

```
In [55]: from sklearn.neighbors import KNeighborsClassifier
ModelKNN = KNeighborsClassifier(n_neighbors=5)
ModelLR.fit(x_train, y_train)

y_pred = ModelLR.predict(x_test)

y_pred_prob = ModelLR.predict_proba(x_test)

print('Model Name: ',ModelLR )

from sklearn.metrics import confusion_matrix
actual = y_test
predicted = y_pred
matrix = confusion_matrix(actual,predicted, labels=[1,0],sample_weight=None, r
print('Confusion matrix : \n', matrix)
```

Model Name: LogisticRegression()  
Confusion matrix :  
[[15534 3094]  
[ 2296 21937]]

```
In [56]: tp, fn, fp, tn = confusion_matrix(actual,predicted,labels=[1,0]).reshape(-1)
print('Outcome values : tp:{} fn:{} fp:{} tn:{}\n'.format(tp, fn, fp, tn) )

sensitivity = round(tp/(tp+fn), 3)
specificity = round(tn/(tn+fp), 3)
accuracy = round((tp+tn)/(tp+fp+tn+fn), 3)
precision = round(tp/(tp+fp), 3)
f1Score = round((2*tp/(2*tp + fp + fn)), 3)

from math import sqrt

print('Accuracy :', round(accuracy*100, 2),'%')
print('Precision :', round(precision*100, 2),'%')
print('Recall :', round(sensitivity*100,2), '%')
print('F1 Score :', f1Score)
print('Specificity or True Negative Rate :', round(specificity*100,2), '%' )
```

Outcome values : tp:15534 fn:3094 fp:2296 tn:21937

Accuracy : 87.4 %  
Precision : 87.1 %  
Recall : 83.4 %  
F1 Score : 0.852  
Specificity or True Negative Rate : 90.5 %

## RandomForestClassifier

```
In [57]: from sklearn.ensemble import RandomForestClassifier
ModelRF = RandomForestClassifier()
ModelRF.fit(x_train, y_train)

y_pred = ModelRF.predict(x_test)

y_pred_prob = ModelRF.predict_proba(x_test)

print('Model Name: ',ModelRF )

from sklearn.metrics import confusion_matrix
actual = y_test
predicted = y_pred
matrix = confusion_matrix(actual,predicted, labels=[1,0],sample_weight=None, r
print('Confusion matrix : \n', matrix)
```

Model Name: RandomForestClassifier()  
Confusion matrix :  
[[17545 1083]  
[ 550 23683]]

```
In [58]: tp, fn, fp, tn = confusion_matrix(actual,predicted,labels=[1,0]).reshape(-1)
print('Outcome values : tp:{} fn:{} fp:{} tn:{}\n'.format(tp, fn, fp, tn) )

sensitivity = round(tp/(tp+fn), 3)
specificity = round(tn/(tn+fp), 3)
accuracy = round((tp+tn)/(tp+fp+tn+fn), 3)
precision = round(tp/(tp+fp), 3)
f1Score = round((2*tp/(2*tp + fp + fn)), 3)

from math import sqrt

print('Accuracy :', round(accuracy*100, 2),'%')
print('Precision :', round(precision*100, 2),'%')
print('Recall :', round(sensitivity*100,2), '%')
print('F1 Score :', f1Score)
print('Specificity or True Negative Rate :', round(specificity*100,2), '%' )
```

Outcome values : tp:17545 fn:1083 fp:550 tn:23683

Accuracy : 96.2 %

Precision : 97.0 %

Recall : 94.2 %

F1 Score : 0.956

Specificity or True Negative Rate : 97.7 %

In [ ]:

In [ ]: